

End-to-End Credit Card Fraud Detection with Machine Learning Models

```
In [3]: # import libraries
import pandas as pd
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings('ignore')
```

```
In [5]: # Load the dataset
df = pd.read_csv(r'C:\Users\User\Downloads\creditcard.csv')
df.head()
```

```
Out[5]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.0986
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.0851
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.2476
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.3774
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.2705

5 rows × 31 columns

```
In [7]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null float64
 1   V1          284807 non-null float64
 2   V2          284807 non-null float64
 3   V3          284807 non-null float64
 4   V4          284807 non-null float64
 5   V5          284807 non-null float64
 6   V6          284807 non-null float64
 7   V7          284807 non-null float64
 8   V8          284807 non-null float64
 9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

This project utilizes the **Credit Card Fraud Detection dataset** from Kaggle, which contains anonymized transactions made by European cardholders in September 2013. The dataset consists of **284,807 transactions** and is characterized by an **extremely imbalanced target variable**, with fraudulent transactions representing approximately **0.17%** of the data.

Metric	Value
Number of rows	284,807
Number of columns	31
Features	28 PCA components (V1–V28), Time, Amount
Target Variable	Class (0 = Non-fraud, 1 = Fraud)
Class Distribution	~0.17% fraud cases (492 frauds)

The dataset includes:

- **Time** — Seconds elapsed between each transaction and the first transaction
 - **V1–V28** — PCA-transformed features (due to confidentiality)
 - **Amount** — Transaction amount
 - **Class** — Binary target variable (fraud indicator)
-

```
In [9]: df.Class.value_counts()
```

```
Out[9]: Class
0      284315
1         492
Name: count, dtype: int64
```

Data is imbalanced, we need to handle the imbalanced data

Problem Statement

Fraudulent credit card transactions result in significant financial losses and customer dissatisfaction for financial institutions.

The goal of this project is to build and benchmark machine learning classification models that can accurately detect fraudulent transactions within a highly imbalanced dataset, while minimizing false negatives and improving fraud detection efficiency.



Project Objective

- To address **class imbalance** using undersampling techniques.
 - To train and evaluate multiple **classification models** (Logistic Regression, KNN, Decision Tree, Random Forest, AdaBoost).
 - To optimize model performance through **hyperparameter tuning** and compare model results.
 - To identify the most effective model(s) for detecting fraudulent credit card transactions.
-

Data Preprocessing

Balancing the data

Undersampling

```
In [425... # dependent and independent data
x = df.drop('Class', axis = 1)
```

```
y = df.Class
```

```
In [427... # train test split
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [429... x_train
```

```
Out[429...
      Time      V1      V2      V3      V4      V5      V6      V
228122 145367.0  2.053311  0.089735 -1.681836  0.454212  0.298310 -0.953526  0.15200
283298 171505.0  2.072657 -0.059176 -1.127889  0.382489 -0.078354 -1.131236  0.17676
49603  44132.0  0.041351 -2.448057 -0.455488  0.406956 -1.569721 -0.870497  0.72670
238026 149493.0  1.990593 -0.330605 -0.920831  0.041564 -0.286964 -0.653376 -0.24149
138828  82866.0  0.762289 -0.301791 -1.368812  1.679729  2.128077  3.782476 -0.31997
...      ...      ...      ...      ...      ...      ...      ...
42487  41121.0 -0.687671  0.162012  0.637712 -0.341793  1.855939  4.173013 -0.66666
161032 113807.0 -0.017840  0.755066  0.613992 -0.647947  0.044703 -0.885837  0.36294
46440  42787.0 -0.341879  0.636188  2.053224  0.600069  0.176559  0.512507  0.45655
211250 138329.0 -0.617552  0.005754  0.001910 -0.777948  0.635097 -0.466782  1.15854
145579  87062.0 -5.068390 -3.517645 -1.442207 -0.503810  7.459502  0.650672 -2.12290
```

199364 rows × 30 columns

```
In [431... y_train
```

```
Out[431...
228122    0
283298    0
49603     0
238026    0
138828    0
..
42487     0
161032    0
46440     0
211250    0
145579    0
Name: Class, Length: 199364, dtype: int64
```

```
In [433... y_train.value_counts()
```

```
Out[433...
Class
0    199009
1       355
Name: count, dtype: int64
```

```
In [435... # undersampling
from imblearn.under_sampling import NearMiss
```

```
In [437... # initialize

nm = NearMiss()
```

```
In [439... x_train_nm, y_train_nm = nm.fit_resample(x_train, y_train)
```

```
In [441... y_train_nm.value_counts()
```

```
Out[441... Class
0      355
1      355
Name: count, dtype: int64
```

```
In [443... x_test_nm, y_test_nm = nm.fit_resample(x_test, y_test)
```

```
In [445... x_test_nm
```

```
Out[445...
      Time      V1      V2      V3      V4      V5      V6      V
0  7551.0  0.861827  1.525916 -3.253113  0.707706  3.561780  2.433556  0.27902
1  7569.0  1.125274  0.715678  1.033160  2.609092 -0.134327 -0.263141 -0.02468
2  7572.0 -0.300229  1.260226  2.276756  1.366404 -0.068287 -0.582341  0.63639
3  93824.0  1.976929 -0.302019 -0.082737  0.739391 -0.767913 -0.582706 -0.69810
4  7567.0 -1.034925  0.421541  1.719689 -0.089020  2.015857  4.770585 -0.65684
...      ...      ...      ...      ...      ...      ...      ...
269  65728.0  1.227614 -0.668974 -0.271785 -0.589440 -0.604795 -0.350285 -0.48636
270  14073.0 -4.153014  8.204797 -15.031714  10.330100 -3.994426 -3.250013 -10.41569
271  155965.0 -1.201398  4.864535 -8.328823  7.652399 -0.167445 -2.767695 -3.17642
272  34521.0  1.081234  0.416414  0.862919  2.520863 -0.005021  0.563341 -0.12337
273  93888.0 -10.040631  6.139183 -12.972972  7.740555 -8.684705 -3.837429 -11.90770
```

274 rows × 30 columns

Model Building

Trained five classifiers:

- Logistic Regression
- K-Nearest Neighbors (KNN)
- Decision Tree Classifier

- Random Forest Classifier
- AdaBoost Classifier

Logistic Regression

```
In [59]: from sklearn.linear_model import LogisticRegression
```

```
In [61]: LR = LogisticRegression()

LR.fit(x_train_nm, y_train_nm)
```

```
Out[61]: LogisticRegression
LogisticRegression()
```

```
In [63]: # Predict on the balanced TEST set
y_pred_test_lr = LR.predict(x_test_nm)
```

```
In [65]: y_pred_test_lr
```

```
Out[65]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
                1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
                1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1,  
                1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
                1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
                1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1,  
                1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1],  
               dtype=int64)
```

```
In [45]: # Evaluation of the Logistic regression

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [67]: print('Accuracy score is: ', accuracy_score(y_test_nm, y_pred_test_lr))
```

Accuracy score is: 0.9401408450704225

```
In [69]: print('Confusion Matrix:\n', confusion_matrix(y_test_nm, y_pred_test_lr))
print('Classification Report:\n', classification_report(y_test_nm, y_pred_test_lr))
```

Confusion Matrix:

```
[[138  4]
 [ 13 129]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.97	0.94	142
1	0.97	0.91	0.94	142
accuracy			0.94	284
macro avg	0.94	0.94	0.94	284
weighted avg	0.94	0.94	0.94	284

In []:

KNN Classifier

In [74]: `from sklearn.neighbors import KNeighborsClassifier`

Initialize and Train the KNN model

In [537... `from sklearn.preprocessing import StandardScaler`

```
scaler = StandardScaler()
x_train_nm_scaled = scaler.fit_transform(x_train_nm)
x_test_nm_scaled = scaler.fit_transform(x_test_nm)
```

In [539... `from sklearn.neighbors import KNeighborsClassifier`

```
knn = KNeighborsClassifier(n_neighbors=3, weights='distance')
knn.fit(x_train_nm_scaled, y_train_nm)
```

Out[539... `KNeighborsClassifier`

```
KNeighborsClassifier(n_neighbors=3, weights='distance')
```

In [541... `y_pred_knn = knn.predict(x_test_nm_scaled)`

In [543... `print("KNN Accuracy:", accuracy_score(y_test_nm, y_pred_knn))`

KNN Accuracy: 0.9343065693430657

Decision Tree Classifier

In [106... `from sklearn.tree import DecisionTreeClassifier`

```
In [110...] dt = DecisionTreeClassifier(random_state=42)
```

```
In [112...] dt.fit(x_train_nm, y_train_nm)
```

```
Out[112...] DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)
```

```
In [114...] y_pred_test_dt = dt.predict(x_test_nm)
```

```
In [116...] from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

print('Accuracy Score:', accuracy_score(y_test_nm, y_pred_test_dt))
print('Confusion Matrix:\n', confusion_matrix(y_test_nm, y_pred_test_dt))
print('Classification Report:\n', classification_report(y_test_nm, y_pred_test_dt))
```

Accuracy Score: 0.6197183098591549

Confusion Matrix:

```
[[ 45  97]
```

```
 [ 11 131]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.32	0.45	142
1	0.57	0.92	0.71	142
accuracy			0.62	284
macro avg	0.69	0.62	0.58	284
weighted avg	0.69	0.62	0.58	284

Hyperparameter tuning

```
In [119...] # Hyperparameter tuning

parameters = {
    'criterion' : ['gini', 'entropy', 'log_loss'],
    'splitter' : ['best', 'random'],
    'max_depth' : [1,2,3,4,5,6,7],
    'max_features' : ["auto", "sqrt", "log2"]
}
```

```
In [121...] # GridSearchCV (CV= Cross Validation)

from sklearn.model_selection import GridSearchCV

treemodel = DecisionTreeClassifier()

gscv = GridSearchCV(treemodel, param_grid=parameters, cv=5, scoring='accuracy')

gscv.fit(x_train, y_train)
```



```
Out[121...] ▶ GridSearchCV
                ▶ best_estimator_:
                  DecisionTreeClassifier
                    ▶ DecisionTreeClassifier
```

```
In [123...] gscv.best_params_
```

```
Out[123...] {'criterion': 'log_loss',
             'max_depth': 7,
             'max_features': 'sqrt',
             'splitter': 'best'}
```

```
In [125...] dt_model = DecisionTreeClassifier(criterion='log_loss', max_depth = 7, max_features
```

```
In [127...] dt_model.fit(x_train_nm, y_train_nm)
```

```
Out[127...] ▼ DecisionTreeClassifier
DecisionTreeClassifier(criterion='log_loss', max_depth=7, max_features='sq
rt')
```

```
In [129...] # Prediction

y_pred_train_DT = dt_model.predict(x_train_nm)

y_pred_test_DT = dt_model.predict(x_test_nm)
```

```
In [131...] print(f'Training Accuracy {accuracy_score(y_pred_train, y_train_nm)}')
print(f'Testing Accuracy {accuracy_score(y_pred_test, y_test_nm)}')
```

Training Accuracy 0.96
Testing Accuracy 0.9119718309859155

Random Forest Classifier

```
In [485...] from sklearn.ensemble import RandomForestClassifier
```

```
In [487...] rf = RandomForestClassifier(random_state=42)
```

```
In [489...] rf.fit(x_train_nm, y_train_nm)
```

Out[489...

```
RandomForestClassifier
RandomForestClassifier(random_state=42)
```

In [491...

```
y_pred_test_rf = rf.predict(x_test_nm)
```

In [493...

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
print('Accuracy Score:', accuracy_score(y_test_nm, y_pred_test_rf))
```

Accuracy Score: 0.9306569343065694

Hyperparameter tuning

In [392...

```
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
```

In [508...

```
params = { 'n_estimators' : [100,200,300],
           'criterion' : ['gini','entropy'],
           'max_depth' : [5,10,15,None],
           'min_samples_split':np.arange(2,10),
           'min_samples_leaf' : np.arange(1,10)
}
```

Grid Search CV

In [510...

```
rf_gscv = GridSearchCV(rf, param_grid= params)
```

In [515...

```
rf_gscv.fit(x_train, y_train)
```

Randomized Search CV

In [400...

```
rfrs = RandomizedSearchCV(rf, param_distributions = params, cv=10, n_iter=15)
```

In [404...

```
rfrs.fit(x_train_nm, y_train_nm)
```

Out[404...

```
RandomizedSearchCV
  best_estimator_:
    RandomForestClassifier
      RandomForestClassifier
```

In [496...

```
rfrs.best_params_
```

```
Out[496... {'n_estimators': 100,
            'min_samples_split': 8,
            'min_samples_leaf': 2,
            'max_depth': 15,
            'criterion': 'gini'}
```

```
In [498... rf_model = RandomForestClassifier(n_estimators = 100, criterion='gini', max_depth =
```

```
In [500... rf_model.fit(x_train_nm, y_train_nm)
```

```
Out[500... ▼ RandomForestClassifier ⓘ ?
RandomForestClassifier(max_depth=15, min_samples_leaf=2, min_samples_split=8)
```

```
In [502... # Prediction
```

```
y_pred_train_RF = rf_model.predict(x_train_nm)
```

```
y_pred_test_RF = rf_model.predict(x_test_nm)
```

```
In [504... print('Accuracy Score (final):', accuracy_score(y_test_nm, y_pred_test_RF))
print('Confusion Matrix:\n', confusion_matrix(y_test_nm, y_pred_test_rf))
print('Classification Report:\n', classification_report(y_test_nm, y_pred_test_rf))
```

Accuracy Score (final)): 0.927007299270073

Confusion Matrix:

```
[[128  9]
 [ 10 127]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.93	0.93	137
1	0.93	0.93	0.93	137
accuracy			0.93	274
macro avg	0.93	0.93	0.93	274
weighted avg	0.93	0.93	0.93	274

AdaBoost Classifier

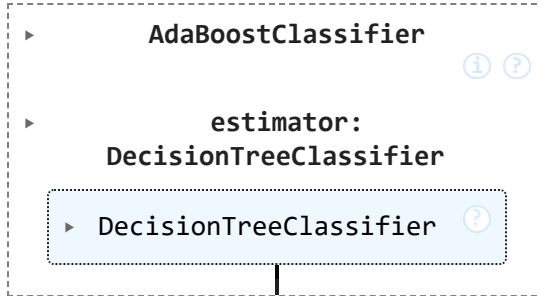
```
In [461... from sklearn.ensemble import AdaBoostClassifier
```

```
In [463... base_est = DecisionTreeClassifier(random_state=42)
```

```
In [465... # Define base estimator (Decision stump)
adaboost = AdaBoostClassifier(estimator=base_est, random_state=42)
```

```
# Initialize AdaBoost with estimator (new versions only)
adaboost.fit(x_train_nm, y_train_nm)
```

Out[465...



In [467...

```
y_pred_test_ada = adaboost.predict(x_test_nm)
```

In [469...

```
## Evaluation parameters
print('Accuracy Score:', accuracy_score(y_test_nm, y_pred_test_ada))
```

Accuracy Score: 0.6496350364963503

Hyperparameter tuning

In [219...

```
# Define parameter grid
params = {
    'n_estimators': [50, 100, 150],
    'learning_rate': [0.01, 0.1, 1],
    'estimator__max_depth': [1, 2, 3]
}
```

In [223...

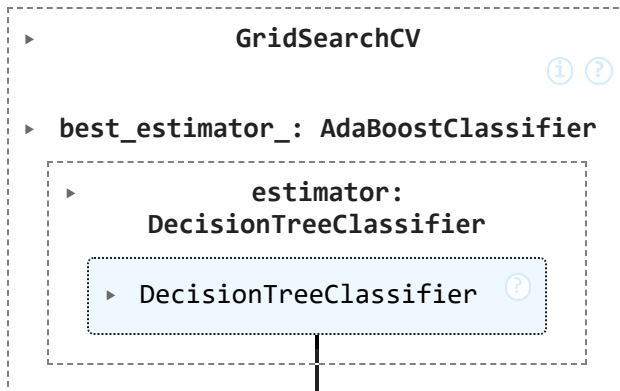
```
treemodel = AdaBoostClassifier()

# GridSearchCV
gscv = GridSearchCV(estimator=adaboost, param_grid=params, cv=5, scoring='accuracy')
```

In [225...

```
gscv.fit(x_train_nm, y_train_nm)
```

Out[225...



In [471...

```
gscv.best_params_
```

Out[471...

```
{'estimator__max_depth': 3, 'learning_rate': 1, 'n_estimators': 100}
```

```
In [473... # Retrain final model
ADAB = AdaBoostClassifier(learning_rate= 1, n_estimators= 100, random_state=42)

ADAB.fit(x_train_nm, y_train_nm)
```

```
Out[473... AdaBoostClassifier

AdaBoostClassifier(learning_rate=1, n_estimators=100, random_state=42)
```

```
In [523... # Predict on test data
y_pred_test_ADAB = ADAB.predict(x_test_nm)
```

```
In [525... # Evaluate

print('Test Accuracy:', accuracy_score(y_test_nm, y_pred_test_ADAB))
print('Confusion Matrix:\n', confusion_matrix(y_test_nm, y_pred_test_ADAB))
print('Classification Report:\n', classification_report(y_test_nm, y_pred_test_ADAB))
```

Test Accuracy: 0.8321167883211679

Confusion Matrix:

```
[[ 98  39]
```

```
[  7 130]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.72	0.81	137
1	0.77	0.95	0.85	137
accuracy			0.83	274
macro avg	0.85	0.83	0.83	274
weighted avg	0.85	0.83	0.83	274

```
In [ ]:
```

Model Comparison and Interpretation

```
In [545... from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt

# Model names
models = ['Logistic Regression', 'KNN Classifier', 'Decision Tree', 'Random Forest']

# Predicted values from each model
y_pred_test_lr = LR.predict(x_test_nm)
y_pred_knn = knn.predict(x_test_nm_scaled)
y_pred_test_DT = dt_model.predict(x_test_nm)
y_pred_test_rf = rf.predict(x_test_nm)
y_pred_test_ADAB = ADAB.predict(x_test_nm)

# Calculate R2 scores
```


```

accuracy_scores = [
    accuracy_score(y_test_nm, y_pred_test_lr),
    accuracy_score(y_test_nm, y_pred_knn),
    accuracy_score(y_test_nm, y_pred_test_DT),
    accuracy_score(y_test_nm, y_pred_test_rf),
    accuracy_score(y_test_nm, y_pred_test_ADAB)
]

# Plot
plt.figure(figsize=(10,5))
plt.bar(models, accuracy_scores, color='cyan', width=0.4)
# Annotate accuracy score values on bars
for i, v in enumerate(accuracy_scores):
    plt.text(i, v + 0.02, f"{v:.2f}", ha='center', fontweight='bold')

plt.xlabel('Models')
plt.ylabel('Accuracy Score')
plt.title('Classification Models Comparison (Accuracy Score)')
plt.ylim(0, 1) # Optional: to make chart scale uniform
plt.show()

```



Logistic Regression and Decision Tree Classifier model gives highest accuracy for this dataset.

Thank you!

In []: